



# **Agenda: Robust Personalized PageRanks in Evolving Graphs**

**Dingheng Mo, Siqiang Luo**

# Personalized PageRank



The random walk interpretation of Personalized PageRank:

- $\pi(s, t)$  = probability that a random walk from  $s$  stops at  $t$

A  $\alpha$ -decay random walk from  $s$

- Stop and restart from  $s$  with  $\alpha$  probability at each step.
- If not stop, randomly jumps to one of its out-neighbors.

# Approximate single-source PPR query

On graph  $G(V, E)$ . With a source node  $s$ , for each node  $v$ , return an estimation value of PPR from  $s$  to  $v$ , denote as  $\hat{\pi}(s, v)$ .

Guarantee the estimation incurs at most  $\epsilon$  relative error

$$|\pi(s, v) - \hat{\pi}(s, v)| \leq \epsilon \cdot \pi(s, v)$$

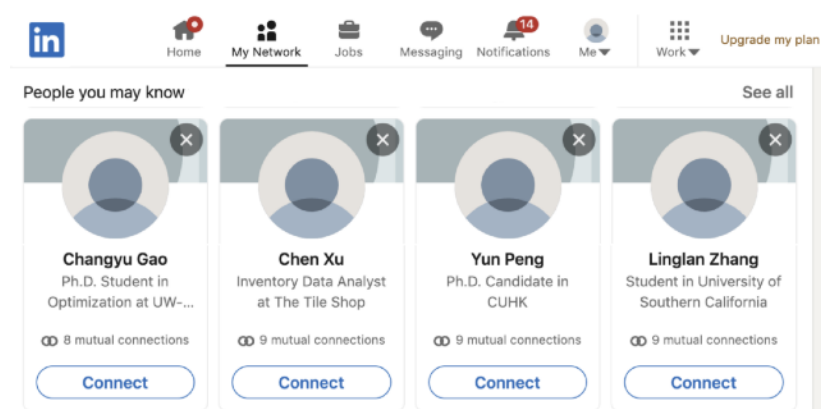
for any  $\pi(s, v) > \delta$ .

# Application

- Personalized search and recommendation




Twitter Who-to-Follow Service



Linkedin user recommendation

# Motivation

Social networks in real applications are highly dynamic.



# Problem definition

With current graph  $G(V, E)$

**Update( $u, v$ ):**


- If edge  $(u, v)$  is not in  $E$ , insert  $(u, v)$  into the graph.
- If edge  $(u, v)$  is already in  $E$ , delete  $(u, v)$  from the graph.

**Query( $s$ ):**

- For each node  $v$ , return the approximate single source PPR value from  $s$  to  $v$  on current graph with accuracy guarantee.

Given a mixture of queries and updates, we aim to design query&update scheme to optimize the total running time of finishing all these queries and updates.

# Previous solution: Monte-carlo



Basic idea: Simulate personalized pagerank by sampling random walks.


Compute  $\pi(s, t)$ : Sample random walks from  $s$ , return fraction of walks end at  $t$ .

[Bahmani VLDB'10] On dynamic graph, index and update random walk segments.

- Requires to store all the intermediate nodes of random walks to maintain the index.
- Stores  $O(\frac{n^2 \log n}{\epsilon^2})$  random walks for queries on whole graph, which is too large.

# Previous solution: Forward push

$\alpha=0.2$



The forward push algorithm maintains a residue and a reserve for each node  $v$

- Reserve : random walks that stops at node  $v$ .
- Residue : random walks that are currently at node  $v$ .

Forward push operation on node  $v$

- Transfers  $\alpha$  portion of its residue to its reserve
- Equally distributes the remaining  $(1 - \alpha)$  portion to the out-neighbors of  $v$ .





# State-of-the-art solution: FORA

The PPR value can be estimated by


$$\pi(s, t) = \underbrace{\pi^{\circ}(s, t)}_{\text{Reserve}} + \sum_{v \in V} \underbrace{r(s, v)}_{\text{Residue}} \cdot \pi(v, t)$$

1. Process forward push from  $s$  to get **reserve** and **residue**.
2. Sample random walks from each node  $v$  to simulate  $\sum_{v \in V} r(s, v) \cdot \pi(v, t)$

# Index



# Challenge



**Index-based** approaches works great for **query heavy** workload.

**Index-free** approaches works great for **update heavy** workload.

There can be **thorny** workloads for both **index-based** and **index-free** approaches.

# Our solution: Agenda


## Agenda

- Is **efficient** on both query and update.
- Maintains index structure incrementally on dynamic graphs.
- Gives low query and update complexity.
- **Faster** than existing index-based and index-free methods on more than **95%** of workloads.
- Up to **100x** improvement than state-of-the-arts on dynamic graphs.

# Intuition

An edge update  $(u, v)$  on graph increase inaccuracy of the random walk index.


However, the increment of inaccuracy is very different on each node.



The index of nodes closer to  $u$  have larger inaccuracy.

# Intuition

Even for the nodes with high index inaccuracy, we still need to consider their importance in query.




# Framework

When it arrives an update( $u, v$ ):

- Perform a reverse push on  $u$  to compute the increment of the index inaccuracy on each node.
- Keep track of the index inaccuracy of each node.

$$\Delta \text{Index inaccuracy}(w) \propto \pi(w, u)$$



# Framework

When it arrives a query(s):

- Perform forward push on  $s$  to get reserve and residue.
- Update indexes on nodes which are overly stale according to index inaccuracy and residue on each node.
- Process random walks to get result with guarantee.



# Low update cost of Agenda

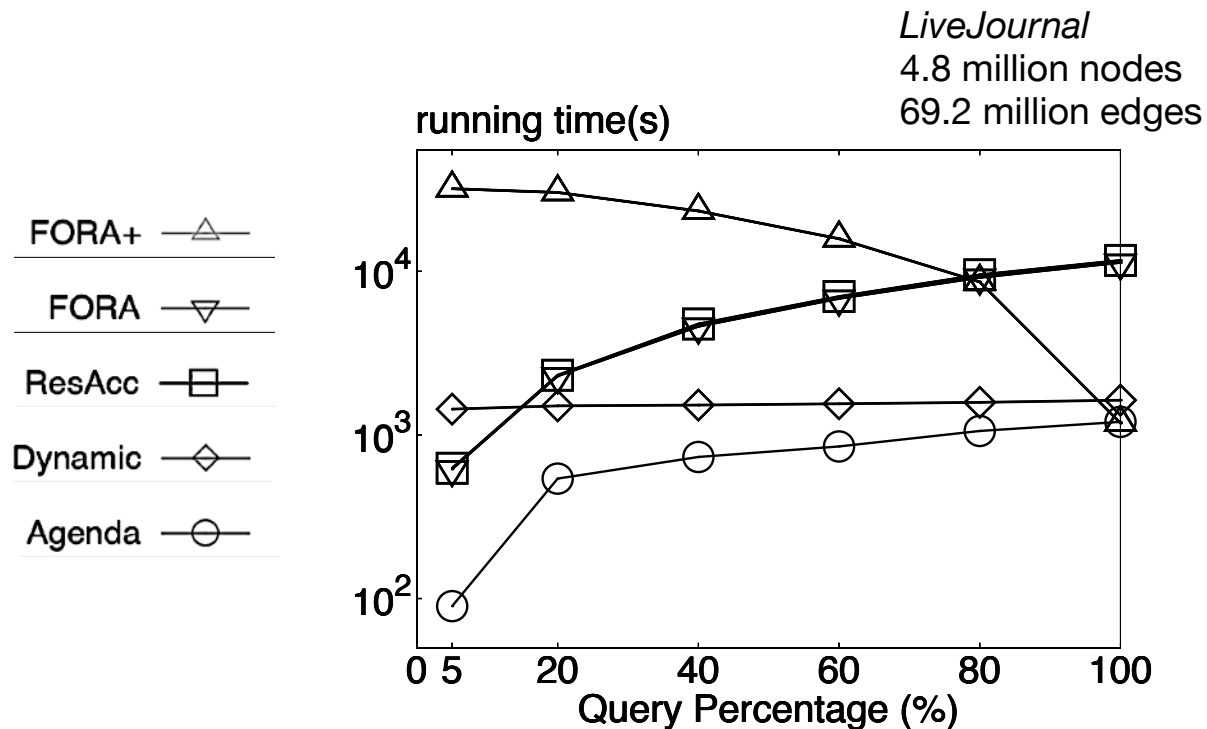
## Agenda guarantees

- In an undirected graph, an edge update incurs at most  $O(\epsilon/\log n)$  fraction of the index to be regenerated.

# Evaluation Baseline

- FORA & FORA+ [Wang et al. KDD'17]
- ResAcc [Lin et al. ICDE'20]
- Dynamic [Bahmani VLDB'10]+[Wang et al. KDD'17]

# Evaluation



Agenda shows better robustness compared with state-of-the-art approaches.

# Conclusion

We present Agenda, an efficient & robust algorithm for single-source PPR queries in dynamic graphs.

## Future work

- Extend our framework to different types of queries.
- Make better trade-off between update, query and space.

Thank you

Q&A